

# Experiment No 11

**Aim:** Using a pre-trained Image net network to predict images into one of the 1000 Imagenet classes

## What is ImageNet?

ImageNet is a large dataset that contains millions of labeled images across 1,000 different categories (e.g., animals, objects, scenes). It has been fundamental in advancing deep learning and computer vision research.

## 2. Pretrained Networks

A pretrained network refers to a neural network model that has already been trained on a specific dataset, such as ImageNet. Popular architectures include ResNet, VGG, and Inception. These models have learned to extract features from images effectively.

## 3. How Prediction Works

When using a pretrained network for image classification, the process generally involves the following steps:

### a. Input Image Preparation

- **Resizing:** The input image is resized to the expected dimensions of the network (often 224x224 pixels for many architectures).
- **Normalization:** Pixel values are typically normalized to a specific range (e.g., 0 to 1 or standardized to have a mean of 0 and a standard deviation of 1).

### b. Feature Extraction

- The image is fed into the pretrained network, which processes it through multiple layers (convolutions, pooling, etc.). Each layer extracts increasingly complex features from the image.

### c. Final Classification Layer

- The output of the final layer is a vector of probabilities corresponding to the 1,000 ImageNet classes. This is usually done using a softmax function, which converts the raw scores into probabilities.

### d. Making Predictions

- The class with the highest probability is selected as the predicted class for the input image.

## 4. Why Use Pretrained Networks?

- **Transfer Learning:** Pretrained models leverage the knowledge gained from training on a large dataset, which helps when you have limited data for your specific task.
- **Reduced Training Time:** You can fine-tune a pretrained model on a smaller dataset, which saves time and computational resources.
- **Improved Performance:** These models often achieve better accuracy than training a new model from scratch.

## 5. Implementation

In practice, using a pretrained network can be done easily with frameworks like TensorFlow or PyTorch. You typically load the model, preprocess your images, and use the model's `predict` function to get predictions.

## 6. Fine-tuning (Optional)

For specific applications, you might want to fine-tune the pretrained model by retraining it on a smaller, domain-specific dataset. This involves:

- Modifying the last layer(s) to match the number of classes in your new dataset.
- Training the model for a few epochs while possibly freezing earlier layers to retain the learned features.

```
In [1]: # !pip install tensorflow==2.1.0
```

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from urllib.request import urlopen
from PIL import Image

%matplotlib inline
plt.style.use('default')

print("TF Version",tf.__version__)
```

TF Version 2.17.0

```
In [3]: model_vgg=tf.keras.applications.vgg16.VGG16(include_top=True, weights='imagenet')
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels.h5)  
553467096/553467096 ————— 4s 0us/step

```
In [4]: model_vgg.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	
input_layer (InputLayer)	(None, 224, 224, 3)	
block1_conv1 (Conv2D)	(None, 224, 224, 64)	
block1_conv2 (Conv2D)	(None, 224, 224, 64)	
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	
block2_conv1 (Conv2D)	(None, 112, 112, 128)	
block2_conv2 (Conv2D)	(None, 112, 112, 128)	
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	
block3_conv1 (Conv2D)	(None, 56, 56, 256)	
block3_conv2 (Conv2D)	(None, 56, 56, 256)	
block3_conv3 (Conv2D)	(None, 56, 56, 256)	
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	
block4_conv1 (Conv2D)	(None, 28, 28, 512)	
block4_conv2 (Conv2D)	(None, 28, 28, 512)	
block4_conv3 (Conv2D)	(None, 28, 28, 512)	
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	
block5_conv1 (Conv2D)	(None, 14, 14, 512)	
block5_conv2 (Conv2D)	(None, 14, 14, 512)	
block5_conv3 (Conv2D)	(None, 14, 14, 512)	
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	
flatten (Flatten)	(None, 25088)	
fc1 (Dense)	(None, 4096)	10
fc2 (Dense)	(None, 4096)	1
predictions (Dense)	(None, 1000)	



Total params: 138,357,544 (527.79 MB)



Trainable params: 138,357,544 (527.79 MB)

Non-trainable params: 0 (0.00 B)

```
In [5]: def preprocess_input(img):
        x=np.zeros((224,224,3),dtype="float32")
        x[:, :,0]=img[:, :,2]
        x[:, :,1]=img[:, :,1]
        x[:, :,2]=img[:, :,0]
        mean = [103.939, 116.779, 123.68]
        x[:, :, 0] = x[:, :, 0]-mean[0]
        x[:, :, 1] = x[:, :, 1]-mean[1]
        x[:, :, 2] = x[:, :, 2]-mean[2]
        return x

        def undo_preprocess_input(img):
            mean = [103.939, 116.779, 123.68]
            img[:, :, 0] = img[:, :, 0]+mean[0]
            img[:, :, 1] = img[:, :, 1]+mean[1]
            img[:, :, 2] = img[:, :, 2]+mean[2]
            x=np.zeros((224,224,3),dtype="float32")
            x[:, :,0]=img[:, :,2]
            x[:, :,1]=img[:, :,1]
            x[:, :,2]=img[:, :,0]
            return x
```



```
In [8]: img1 = (Image.open(urlopen("https://raw.githubusercontent.com/tensorchiefs/dl_book/
plt.imshow(img1)
plt.show()
new_width = 224
new_height = 224
# Replacing Image.ANTIALIAS with Image.Resampling.LANCZOS
img1 = img1.resize((new_width, new_height), Image.Resampling.LANCZOS)
plt.imshow(img1)
plt.show()
img1=np.array(img1)
```



```
In [10]: img2 = (Image.open(urlopen("https://raw.githubusercontent.com/tensorchiefs/dl_book/
plt.imshow(img2)
plt.show()
new_width = 224
new_height = 224
```



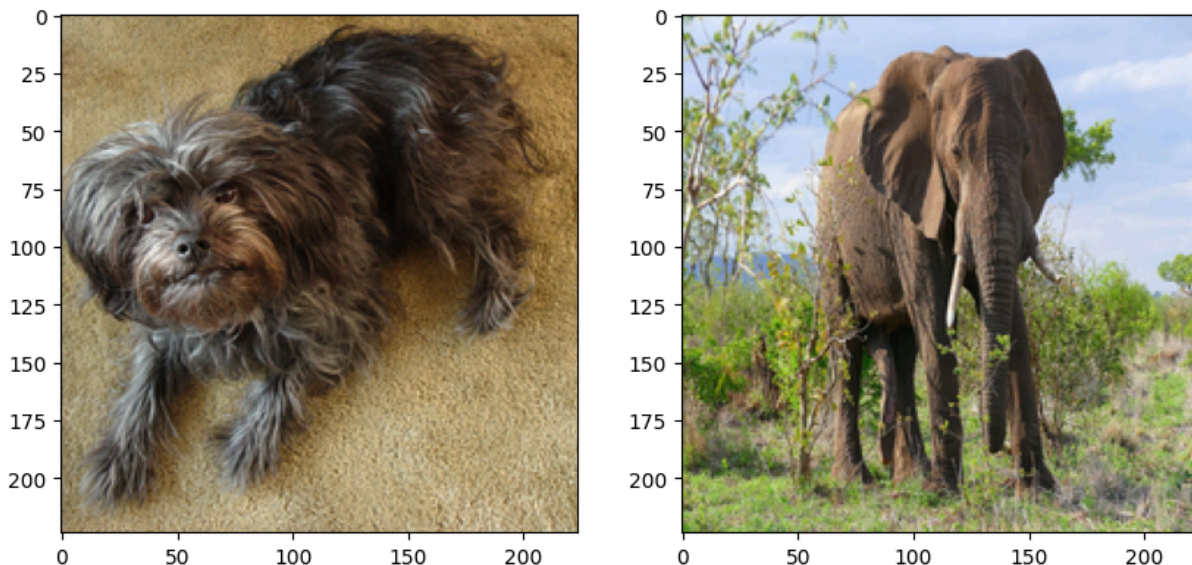
```
# Replacing Image.ANTIALIAS with Image.Resampling.LANCZOS
img2 = img2.resize((new_width, new_height), Image.Resampling.LANCZOS)
plt.imshow(img2)
plt.show()
img2=np.array(img2)
```



```
In [11]: plt.figure(figsize=(10,10))
plt.subplot(1,2,1)
```

```
plt.imshow(img1)
plt.subplot(1,2,2)
plt.imshow(img2)
```

Out[11]: <matplotlib.image.AxesImage at 0x7e0130197a00>



```
In [12]: img1=preprocess_input(img1)
print(img1.shape)
img2=preprocess_input(img2)
print(img2.shape)
```

(224, 224, 3)

(224, 224, 3)

```
In [13]: pred1=model_vgg.predict(np.expand_dims(img1,axis=0))
tf.keras.applications.vgg16.decode_predictions(pred1)
```

1/1 ————— 1s 947ms/step

Downloading data from [https://storage.googleapis.com/download.tensorflow.org/data/imagenet\\_class\\_index.json](https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json)

35363/35363 ————— 0s 0us/step

```
Out[13]: [('n02110627', 'affenpinscher', 0.60996807),
          ('n02094433', 'Yorkshire_terrier', 0.115058616),
          ('n02097474', 'Tibetan_terrier', 0.06349397),
          ('n02086240', 'Shih-Tzu', 0.053632747),
          ('n02098413', 'Lhasa', 0.03696351)]
```

```
In [14]: pred2=model_vgg.predict(np.expand_dims(img2,axis=0))
tf.keras.applications.vgg16.decode_predictions(pred2)
```

1/1 ————— 1s 539ms/step

```
Out[14]: [('n01871265', 'tusker', 0.7851958),
          ('n02504458', 'African_elephant', 0.15363708),
          ('n02504013', 'Indian_elephant', 0.061132655),
          ('n02437312', 'Arabian_camel', 6.303232e-06),
          ('n02109047', 'Great_Dane', 5.9501795e-06)]
```